

Declarative Concurrent Programming with Join Calculus in Scala

Sergei Winitzki

Lambda-Conf 2020 Global Edition

August 11, 2020

Chemical Machine: a new hope

...and some new hype

In this talk:


- “join calculus” as a “**Chemical Machine**”, without academic jargon
 - ▶ “ λ -calculus” — a small (but complete) programming language
 - ▶ “join calculus” — a small (but complete) language for concurrency
 - ★ adds 2 features (“async mailboxes” and “async functions”) to λ -calculus
 - ★ in the chemical metaphor: “molecules” and “reactions”
- Chymyst – open-source implementation of Chemical Machine (Scala)
- examples of concurrent programs in Chymyst
 - ▶ implement anything in 10-15 lines of code
- comparisons with Actor Model, “AWS Lambda”, and Petri nets
- an extension for distributed programming: DCM

Not in this talk: ~~academic theory~~

- ~~Petri net theory, π -calculus, join calculus, joinads, formal semantics~~
- ~~DCM defined within some formalism for distributed programming~~

The chemical metaphor I. “Abstract” chemistry


Real chemistry is asynchronous, concurrent, and parallel:



Want to run *computations* similarly to how chemical reactions run!
Begin by formulating the execution model of “abstract chemistry”:

- Abstract “molecules” float around in a “chemical reaction site”
- Certain kinds of molecules may combine to start a “reaction”:

“abstract” chemical laws:




- Program code defines molecules a , b , c , ... and chemical reactions
- At initial time, the code emits some molecules into the site
- The runtime system runs reactions *concurrently and in parallel*
 - ▶ A chemical simulation engine is easy to implement in any language

The chemical metaphor II. Chemical Machine in a nutshell

From abstract chemistry to computation

Translating the chemical metaphor into a model of computation:

- Each molecule carries a **value** (“concurrent data”)
- Each reaction computes new values from its input values
- Some molecules with new values may be emitted back into the reaction site



```
site(  
  go { case a(x) + b(y) =>  
    val z = f(x, y); a(z) },  
  go { case a(x) + c(_) =>  
    println(x) }  
)
```

When a reaction starts: input molecules disappear, new values are computed, output molecules are emitted

Reactions are *functions* from input values to output values

- Need to learn how to “think in molecules and reactions”

Programming the Chemical Machine using Chymyst

How I learned to forget semaphores and to love concurrency

Molecule emitters are values of type `M[A]`:

```
val count = m[Int]; val inc = m[Unit]
```

Molecules are emitted by calling the emitter's `apply()` method:

```
count(0); inc(); inc(); inc();
```

Reactions are values of type `Reaction`:

```
val r: Reaction = go { case count(x) + inc(_) => count(x + 1) }
```

- A reaction may be any partial function with a single `case` clause
- Emitters (e.g., `count`, `inc`) must be already defined

Reactions are “activated” as a group, using a “reaction site”:

```
site(r1, r2, r3)
```

- Molecules may be emitted after activating their consuming reactions
- A site *must* include all reactions that contend on an input molecule

Programs declare molecules and reaction sites, then emit initial molecules

- The Chymyst runtime runs the simulation

Example: throttling

Throttle emitting molecules `s(x)` with minimum allowed delay of `delta` ms

```
def throttle[X](s: M[X], delta: Long): M[X] = {  
  val r = m[X]  
  val enable = m[Unit] // Molecule is confined to the local scope.  
  site(  
    go { case r(x) + enable(_) =>  
          s(x)  
          Thread.sleep(delta)  
          enable()  
        }  
  )  
  enable() // Enable emitting 's' initially.  
  r // Outside scope will be able to emit 'r'.  
}
```

- No threads/semaphores/locks, no mutable state
- External code may emit `r(x)` at will, and `s(x)` is then throttled
- External code may not emit `enable()`

Implementations in Akka, in Monix, and ZIO: > 50 LOC each

Example: cancellation

Cancel emitting molecules `x(a)` after a molecule `cancel()` is emitted

```
def cancellation[A](x: M[A]): (M[A], M[Unit]) = {  
  val cx = m[A] // Cancellable version of 'x'.  
  val cancel = m[Unit]; val enable = m[Unit]  
  site(  
    go { case cx(a) + enable(_) => x(a); enable() },  
    go { case enable(_) + cancel(_) => }  
  )  
  enable() // Enable emitting 'x' initially.  
  (cx, cancel) // Outside scope will be able to emit these.  
}
```

- No threads/semaphores/locks, no mutable state
- External code may not emit `enable()` because of local scope

Users cannot implement this in Monix or ZIO (cancellation is provided)

An additional feature: Blocking molecules

A molecule emitter type `B[A, R]` receives a reply value of type `R`:

```
val fetch = b[Unit, Int] // Blocking emitter.  
val result = m[Int]  
site(  
  go { case fetch(_, reply) + res(b) => reply(b) }  
) // 'reply' is a ReplyEmitter[Int] defined within reaction scope.  
result(123)  
val x = fetch() // Blocking call, will set x = 123.
```

Blocking emitters are a convenience, do not increase expressive power

Example: map/reduce

A simple map/reduce implementation:

```
val c = m[A] // Initial values have type 'A'.
val d = m[(Int, B)] // 'B' is a commutative monoid.
val res = m[B] // Final result of type 'B'.
val fetch = b[Unit, B] // Blocking emitter.
site(
  // 'map'
  go { case c(x) => d((1, long_computation(x))) },
  // 'reduce'
  go { case d((n1, b1)) + d((n2, b2)) =>
    val (newN, newB) = (n1 + n2, b1 |+| b2)
    if (newN == total) res(newB) else d((newN, newB))
  },
  go { case fetch(_, reply) + res(b) => reply(b) }
)
(1 to 100).foreach(x => c(x))
fetch() // Blocking call will return the final result.
```

Compare with the [Akka implementation here](#) (100+ LOC)

Example: parallel merge-sort

Reactions can be recursive

Complete Chymyst code: [MergeSortSpec.scala](#)


```
val mergesort = m[(Array[T], M[Array[T]])]
site(
  go { case mergesort((arr, sortedResult)) =>
    if (arr.length <= 1) sortedResult(arr)
    else {
      val sorted1 = m[Array[T]]
      val sorted2 = m[Array[T]]
      site( // Define a lower-level reaction site.
        go { case sorted1(x) + sorted2(y) =>
          sortedResult(arrayMerge(x,y))
        }
      )
      val (part1, part2) = arr.splitAt(arr.length/2)
      // Emit lower-level mergesort molecules:
      mergesort(part1, sorted1) + mergesort(part2, sorted2)
    }
  })
```

Implementation in Akka: 30 LOC for the same functionality

Dining philosophers I. Declarative vs. non-declarative code

The paradigmatic example of concurrency, parallelism and resource contention

Five philosophers sit at a round table, taking turns eating and thinking for random time intervals



Problem: simulate the process, avoiding deadlock and starvation

Solutions in various programming languages: see [Rosetta Code](#)

- The Chemical Machine code is purely declarative

Dining philosophers II. Implementation in Chymyst

Five Dining Philosophers implemented in 15 lines of code

Philosophers 1, 2, 3, 4, 5 and forks f12, f23, f34, f45, f51

```
// ... definitions of emitters, think(), eat() omitted for brevity
site (
  go { case t1(_) => think(1); h1() },
  go { case t2(_) => think(2); h2() },
  go { case t3(_) => think(3); h3() },
  go { case t4(_) => think(4); h4() },
  go { case t5(_) => think(5); h5() },

  go { case h1(_) + f12(_) + f51(_) => eat(1); t1() + f12() + f51() },
  go { case h2(_) + f23(_) + f12(_) => eat(2); t2() + f23() + f12() },
  go { case h3(_) + f34(_) + f23(_) => eat(3); t3() + f34() + f23() },
  go { case h4(_) + f45(_) + f34(_) => eat(4); t4() + f45() + f34() },
  go { case h5(_) + f51(_) + f45(_) => eat(5); t5() + f51() + f45() }
)
t1() + t2() + t3() + t4() + t5()

f12() + f23() + f34() + f45() + f51()
```

Source code: [DiningPhilosophers.scala](#)

For more examples, see the [code repository](#) (first-of, barriers, rendezvous, critical sections, readers/writers, Game of Life, elevators, etc.)

Reasoning about code in the Chemical Machine paradigm

Reasoning about concurrent data:

- Emit molecule with value \approx lift data into the “concurrent world”
- Define reaction \approx lift a function into the “concurrent world”
- Reaction site \approx container for concurrent functions and data
- Reaction consumes molecules \approx function consumes input values
- Reaction emits molecules \approx function returns result values

Reasoning about code:


- What data do we need to handle concurrently? (Put it on molecules.)
- What computations consume this data? (Define as reactions.)

Guarantees:

- Molecule emitters and reactions are immutable values in local scopes
- Reaction sites are immutable once activated; can refactor to libraries
- Multiple input molecules are consumed atomically by reactions

Chemical Machine paradigm to become mainstream in 2033

- The gap from academic invention to industry adoption is 38.2 years (infix math, continuations, λ -functions, OOP, CSP, map/reduce, Actor Model, constraint programming, DAG dataflow, Hindley-Milner types)



- The Chemical Machine paradigm was invented in 1995

Current features of Chymyst

Experience with Chemical Machine programming is limited

Some features promise to be useful:

- Blocking molecules with timeouts and timeout back-signalling
- Automatic pipelining of molecules (ordered mailboxes)
- Thread pools, thread priority control, graceful shutdown
- Facility to increase parallelism when using blocking code
- Errors in DSL are reported at compile-time or early run-time
 - ▶ Chymyst uses Scala macros – but only to inspect code
 - ▶ Static analysis enables optimization and error reporting
- “Static” molecules with read-only access (similar to Akka “agents”)
 - ▶ Forgetting to emit a molecule is #1 programmer error
- Logging, debugging, unit-testing facilities

Related frameworks: Petri nets

Workflow management: an approach based on **Petri nets**

- **ING Baker** – a DSL for workflow management, based on Petri nets
- Process modeling and control (“elevator system” etc.)
- Business process management (BPM) systems

Chymyst implements a feature-rich version of Petri nets:

- Transitions admit arbitrary guard conditions and error recovery
- Transitions carry values, reactions are values, can be nested
- Nondeterministic, asynchronous, parallel execution

Any Petri net model is straightforwardly translated into a CM program

Chemical Machine vs. Amazon AWS Lambda

How AWS λ works:

- wait for an event that signals arrival of input data
- run a computation whenever input data becomes available
- the computation is automatically parallelized, data-driven
- writing the output data will create a new event

Modify the AWS λ execution model by adding new requirements:

- a Lambda should be able to wait for several *unrelated* events
- several Lambdas may contend *atomically* on shared input events

With these new requirements, AWS λ becomes “AWS π ” – a model of unrestricted concurrency

- (Implementation on AWS could be tricky)

Chemical Machine vs. the Actor model. I

Modify the Actor execution model by adding new requirements:

- when messages arrive, actors are auto-created, maybe *in parallel*
- actors may wait atomically for messages in *several* different mailboxes

It follows from these requirements that...

- Auto-created actor instances are *stateless* and invisible to user
- User code defines *mailboxes* and *computations* that consume messages
- Repeated messages may be consumed in parallel
- Messages are sent to mailboxes, not to specific actor instances:

<pre>// Akka val a: ActorRef = ... receive(x) =>... val b: ActorRef = ... receive(y) =>... a ! 100 b ! 1; b ! 2; b ! 3</pre>	<pre>// Chymyst ... go { case a(x) => ... } ... go { case b(y) + c(z) => ... } a(100) b(1); b(2); b(3); c("hello");</pre>
--	---

- All data resides on messages in mailboxes, is consumed automatically
- Mailboxes and computations are *values*, can be sent on messages

Any Actor program can be straightforwardly translated into CM

Chemical Machine vs. Actor model. II

- reaction \approx function body for an (auto-started) actor
- emitted molecule with value \approx message with value, in a mailbox
- molecule emitters \approx mailbox references

Programming with actors:

- user code creates and manages explicit actor instances
- actors typically hold mutable state and/or mutate “behavior”
 - ▶ reasoning is about running processes *and* the data sent on messages

Programming with the Chemical Machine:

- processes auto-start when the needed input molecules are available
- many reactions may start at once, with automatic parallelism
 - ▶ user code does not manipulate references to processes
 - ★ no state, no supervision, no lifecycle, no “dead letters”, no routers
 - ▶ reasoning is only about the *data currently available* on molecules
 - ★ no reasoning about running processes having internal state

Chymyst code is typically 2x – 3x shorter than equivalent Akka code

Distributed Chemical Machine

Run concurrent code on a cluster with no code changes

- Declare some molecules as “distributed”, of type `DM[T]`
- No other new language constructions are necessary!
 - ▶ early prototype in progress, as extension of Chymyst

Distributed map/reduce in 15 LOC:

```
implicit val cluster = ClusterConfig(???)
val c = dm[Int] ; val d = dm[Int] // distributed
val res = m[(Int, List[Int])] // local
val fetch = b[Unit, List[Int]]
site(
  go { case c(x) => d(x * 2) }, // ‘map’ on cluster,
  // ‘reduce’ on the driver node only.
  go { case res((n, list)) + d(x) => res((n-1, s::list)) },
  // fetch results
  go { case fetch(_, reply) + res((0, list)) => reply(list) }
)
if (isDriver) { // ‘true’ only on the driver node.
  Seq(1, 2, 3).foreach(x => c(x))
  res((3, Nil)) ; fetch() // Returns the result.
}
```

Comparison: Akka implementation of distributed map/reduce (400+ LOC)

Distributed cache in 10 LOC

- Mutable `Map[String, String]` with operations: `put`, `get`, `delete`

```
implicit val cluster = ClusterConfig(???)
val data = dm[mutable.Map[String, String]]
val put = dm[(String, String)]
val get = dm[(String, M[Option[String]])]
val delete = dm[String]
site(
  go { case data(dict) + put((k, v)) => data(dict.updated(k, v)) },
  go { case data(dict) + get((k, r)) => data(dict); r(dict.get(k)) },
  go { case data(dict) + delete(k) => dict.remove(k); data(dict) }
)
if (isDriver) data(mutable.Map[String, String]())
```

- Comparison: Distributed cache in 100 lines of Akka

Distributed peer-to-peer chat in 15 LOC

- Register user names in chat room
- Fetch list of users
- Send and receive text messages

```
implicit val cluster = ClusterConfig(???)
val users = dm[List[DM[String]]] // List of users' message emitters.
val carrier = dm[DM[String]] // Carries this node's message emitter.
val fetch = b[Unit, List[DM[String]]]
site(go { case users(es) + carrier(e) => users(e :: es) }
, go { case users(es) + fetch(_, r) => users(es); r(es) } )

val peerName = ??? // Read from config on node.
val sender = new DM[String](peerName) // Assign unique molecule name.
site(go { case sender(x) => println(s"Peer $peerName reads $x")})

carrier(sender)
if (isDriver) users(nil)
// Fetch list of users and send a message to Sergei if present.
fetch()
  .find(_.name == "Sergei")
  .foreach(sender => sender("hello"))
```

- Comparison: Distributed chat in > 100 lines of Akka

Reasoning in the Distributed Chemical Machine

Distributed computing is made declarative

- Determine which data needs to be distributed and/or concurrent
- Determine which computations will need to consume that data
- Emit initial molecules and let the DCM run

Pure peer-to-peer architecture:

- Distributed molecules may be consumed by *any* DCM peer
- All DCM peers operate in the same way (no master/worker)
- All DCM peers need to define the same distributed reaction sites
 - ▶ Non-DCM code may differ between peers
 - ▶ Code or configuration could designate DCM peer as a “driver” or have different roles

Chemical Machine: implementation details

- Each reaction site has a scheduler thread and a worker thread pool
- Each molecule is “bound” to a unique reaction site
- Each emitted molecule is stored in a multi-set at its reaction site
- Each emitted molecule triggers a search for possible reactions
 - ▶ Reaction search proceeds concurrently for different reaction sites
- Reactions are scheduled on the worker thread pool
 - ▶ The thread pool can be configured per-reaction or per-site
- Scala macros are used for static analysis and optimizations
 - ▶ Automatically pipelined molecules
 - ▶ Simplify and analyze Boolean conditions
- Error analysis is also performed at early run time
 - ▶ Reaction site with errors remain inactive

Distributed Chemical Machine: implementation details

- Each distributed molecule (DM) is bound to a unique reaction site
- Emitted DM data goes into the ZK instance
- Each DCM peer listens to ZK messages and checks for its DMs
 - ▶ Once a DM is found, its data is downloaded and deserialized
- On a DCM peer, each DM is identified with a unique local RS
 - ▶ Downloaded molecules are emitted into the local RS to run reactions
 - ▶ All DCM peers must run identical reaction code for DMs
- Each DCM peer acquires a distributed lock on its DMs
 - ▶ Lock is released once reaction scheduling is complete
- If a node goes down or network fails, molecules will be *unconsumed*
 - ▶ Another DCM peer will pick up these molecules later

Conclusions and outlook

- Chemical Machine = declarative, purely functional concurrency
 - ▶ Enough power to replace threads, semaphors, atomic vars, etc.
 - ▶ Similar to Actor Model, but easier to use and “more purely functional”
 - ▶ Significantly shorter code, easier to reason about
- An open-source Scala implementation: **Chymyst**
 - ▶ Static DSL code analysis (with Scala macros)
 - ▶ Industry-strength features (thread priority control, pipelining, fault tolerance, unit testing and debugging APIs)
 - ▶ Extensive documentation: **tutorial book**
 - ▶ Distributed Chemical Machine – work in progress
- Promising applications:
 - ▶ Workflow management, BPM
 - ▶ Asynchronous GUIs
 - ▶ Distributed peer-to-peer systems